

# Lecture 12 (Bonus track): Deep Learning for point clouds and graphs

## Motivation (overview)

- Autonomous robotics is a huge frontier for AI (maybe the most impactful one?)
- Early successes: autonomous driving prototypes, delivery bots, vacuum cleaners, warehouse robots...
- 3D sensors = big boost
- Deep learning = big boost
- Deep learning + 3D data =  ?

# Why point clouds are important?



RealSense

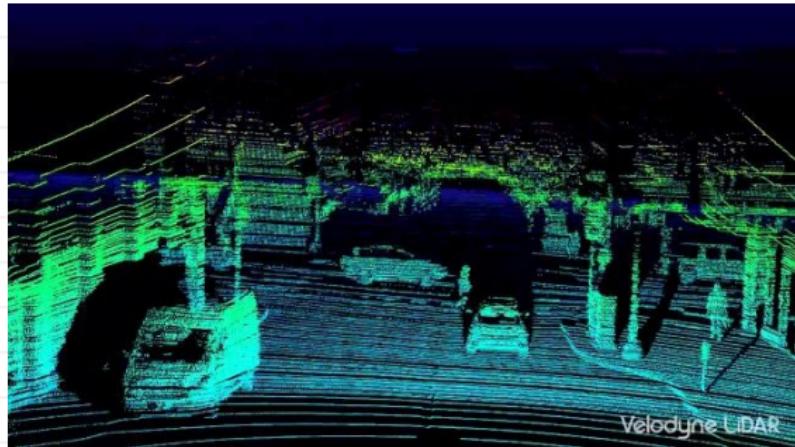


Kinect

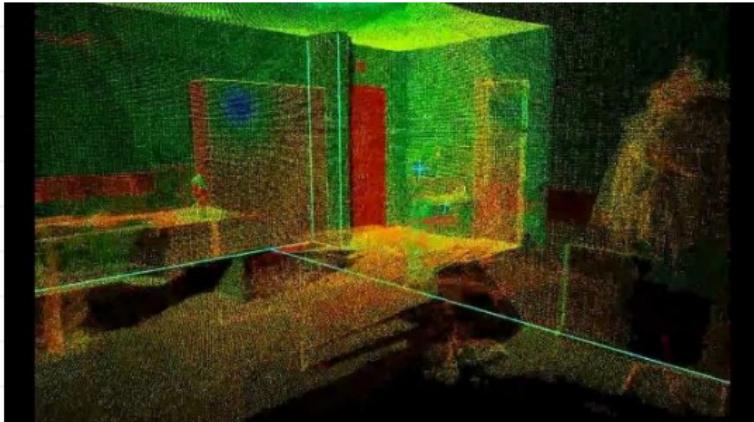


Apple scanner

# Why point clouds are important?



- Dominant sensor for autonomous driving
- Widely used in architecture and urban planning



# Why point clouds are important?

Search for antique armchair ~ 3D Warehouse

3D Warehouse

antique armchair

Upload Model

Sign In

28 Results ALL Results Per Page  Models  Collections Sort by Relevance

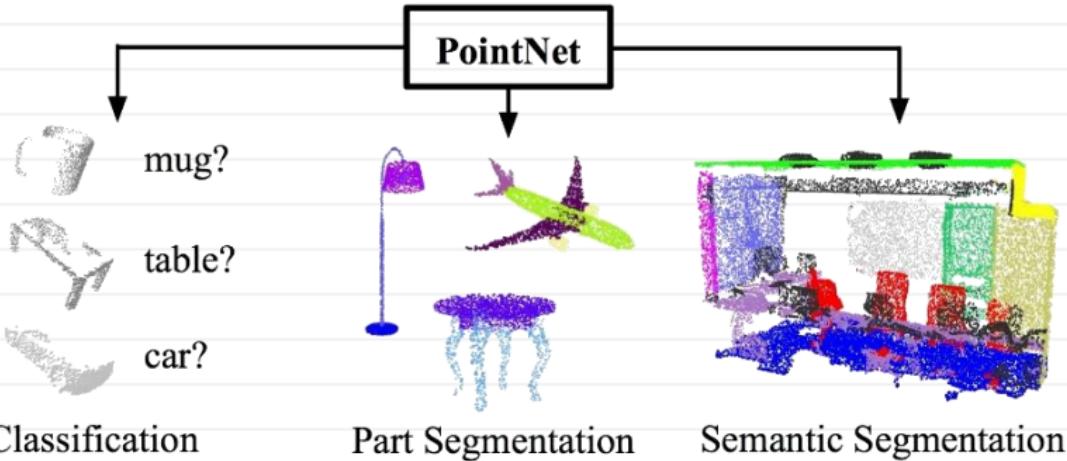
The screenshot shows a grid of 16 thumbnail images of various antique armchairs, each with a download icon. The models are arranged in four rows of four. The titles and creators of the models are listed below their respective thumbnails.

- Gothic Revival Sofa and Armchair by: cealupu
- Stuart Era Armchair by: cealupu
- Hutten 1160/46 by: jvr
- Jacobean Revival Dining Armchair by: cealupu
- Charles I Oak Wainscot Chair by: cealupu
- Hutten CL-20/71 + 162/01+116... by: jvr
- Sleeping Chair 1675 by: cealupu
- Louis XVI Fauteuil, Circa 1780 by: D. William Larson Interior Design
- Antique wooden armchair by: Arthas
- Wooden Arm Chair by: al\_xv
- French President Office , Salon ... by: ADRIEN23
- Hutten 850/71 + 1162/01 by: jvr
- (empty)
- (empty)
- (empty)
- (empty)

Trimble. ©2017 Trimble Inc. Privacy Terms of Use English

Done

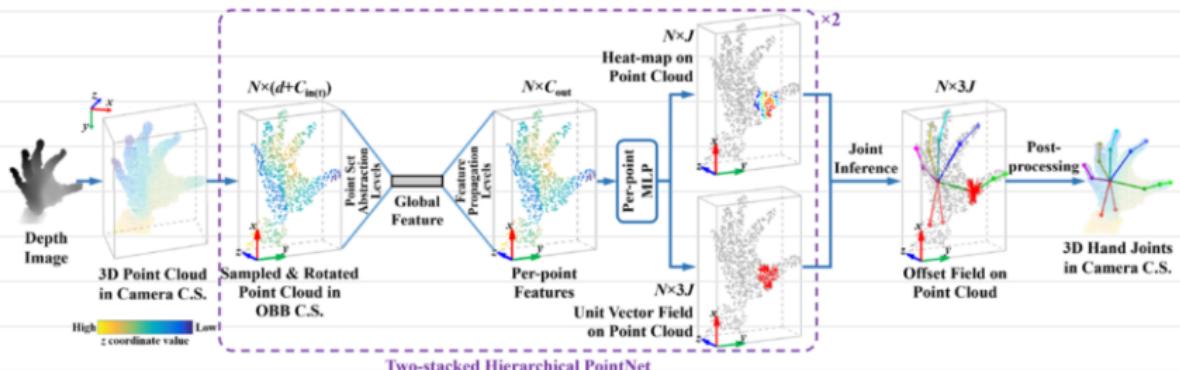
# Recognition tasks



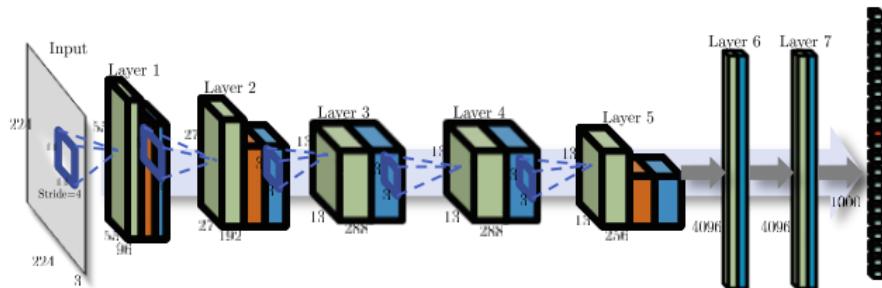
Classification

Part Segmentation

Semantic Segmentation

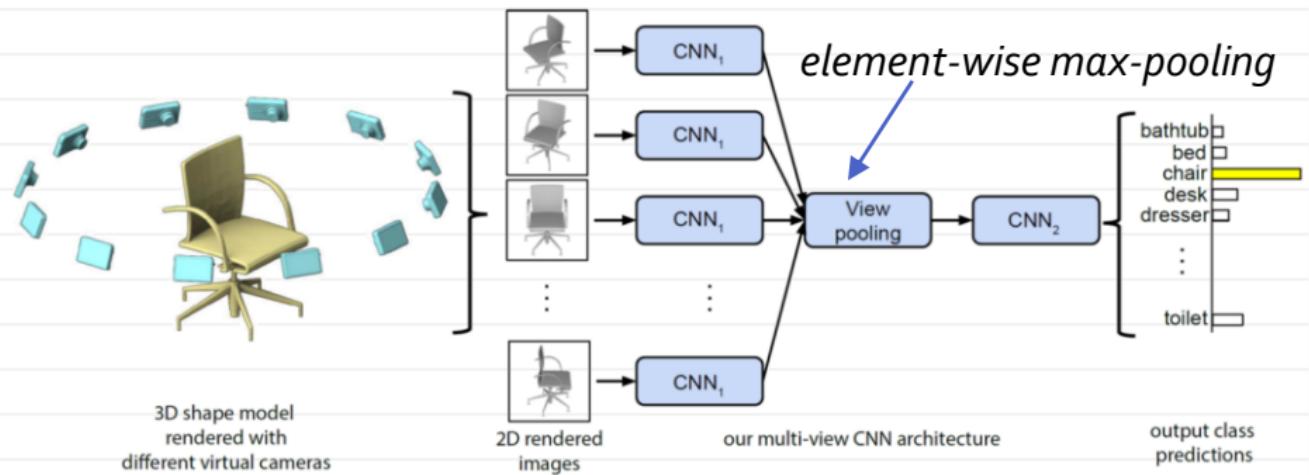


# General motivation



This is ConvNet  
It does recognition very well  
Be like ConvNet

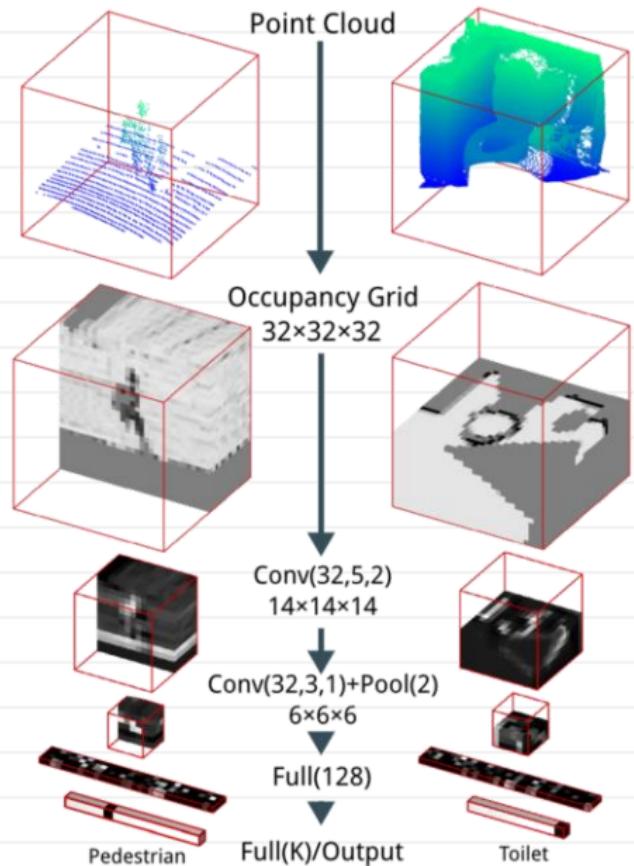
# Multi-view ConvNets



- Take an object (point cloud or mesh)
- Render it to 2D images
- Apply ConvNets

[Su et al. CVPR15]

# 3D volumetric CNNs

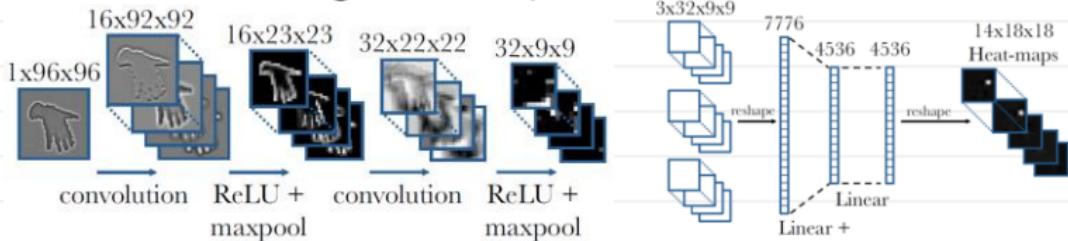


- Straightforward generalization to 3D
- Resolution is very limited

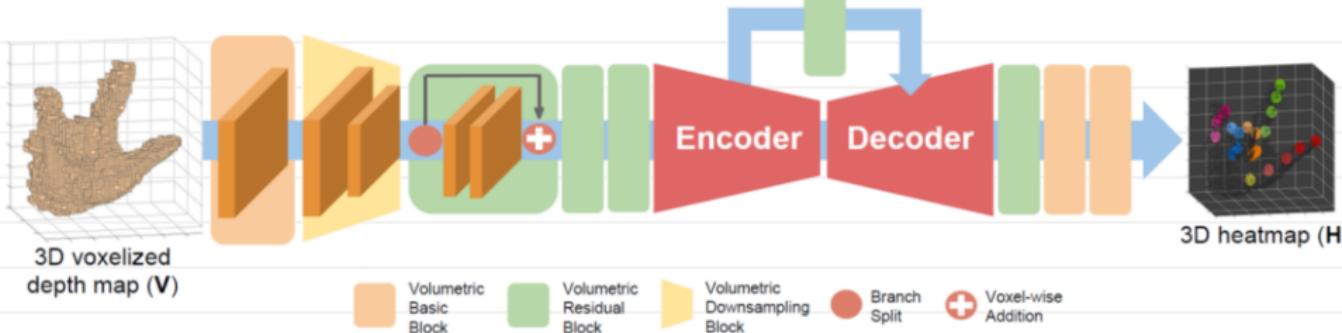
[Maturana IROS15]

# V2V hand tracking

- Before RGBD->hand was handled by 2D ConvNets, e.g. [Thompson SIGGRAPH14]:



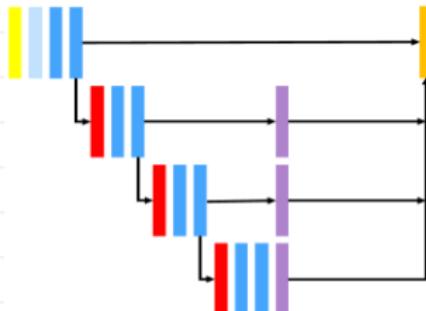
- [Moon et al. CVPR18] showed that 3D ConvNets were better (even with limited time budget)



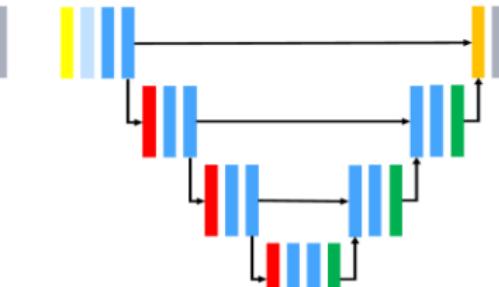
# Submanifold-Sparse ConvNets



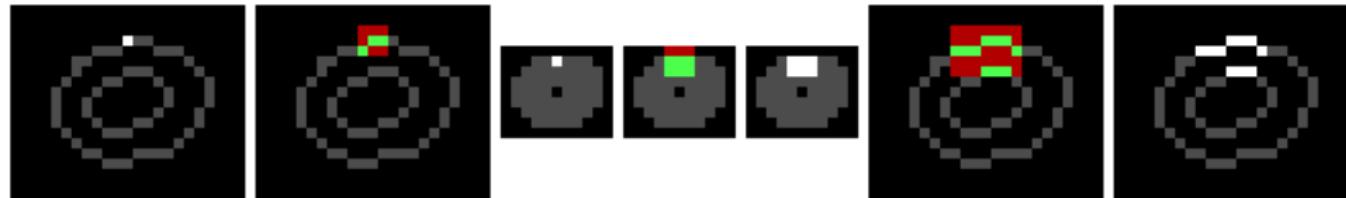
Idea: use sparse format, keep only voxels that were initially occupied



(a) Submanifold sparse FCN.



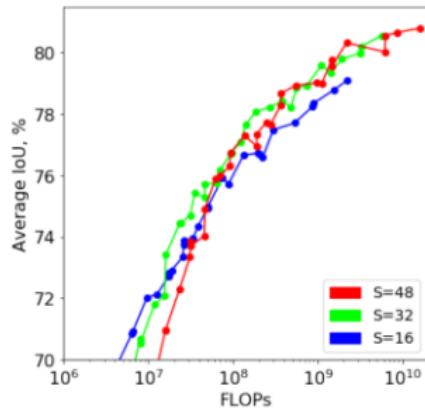
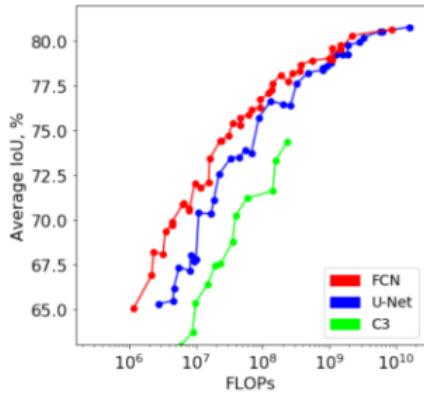
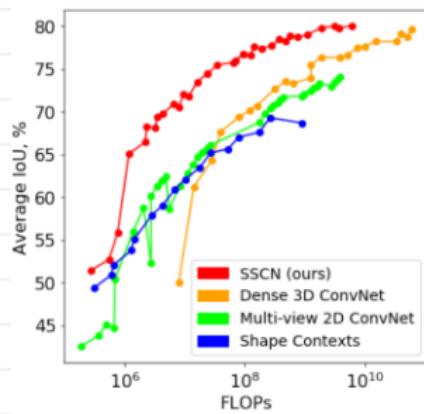
(b) Submanifold sparse U-Net.



[Graham et al. CVPR18]

# Submanifold-Sparse ConvNets

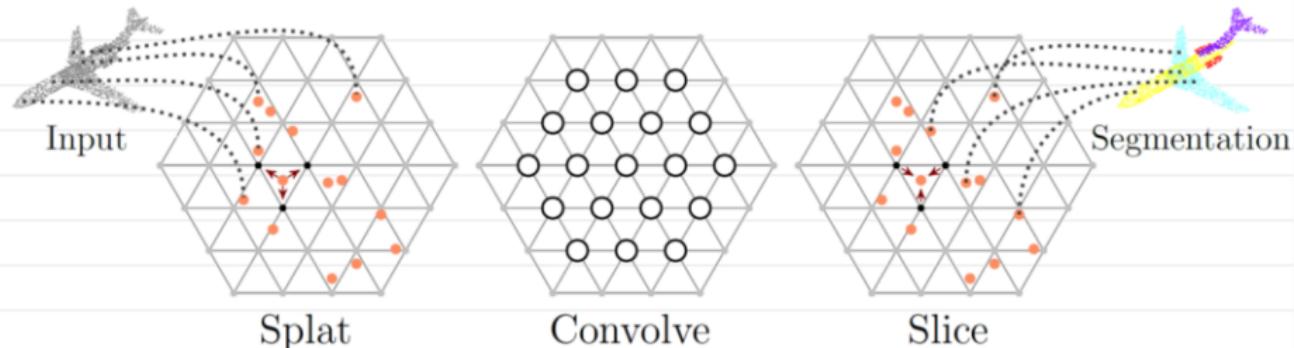
ShapeNet benchmarks:  
(note: FLOPs != runtime)



Other related idea: use octrees to store data (much more tricky)

[Graham et al. CVPR18]

# SplatNet



$$\hat{F}_c = S_{slice} B_{conv} S_{splat} F_c$$

- Permutohedral lattices are more efficient than rectangular grids in high dimensions (not really used in this work)
- SOTA performance (last year) across number of tasks

[Su et al. CVPR18]

# Kd-trees: recap

```
function node = BuildKdTree(M)
```

```
if |M| > 1
```

```
a = axis with the biggest variance
```

```
M = sort(M, axis = a)
```

```
node.child1 = BuildKdTree(M[1:|M|/2])
```

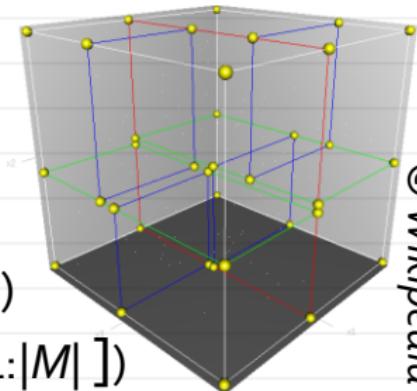
```
node.child2 = BuildKdTree(M[|M|/2+1:|M|])
```

```
return node
```

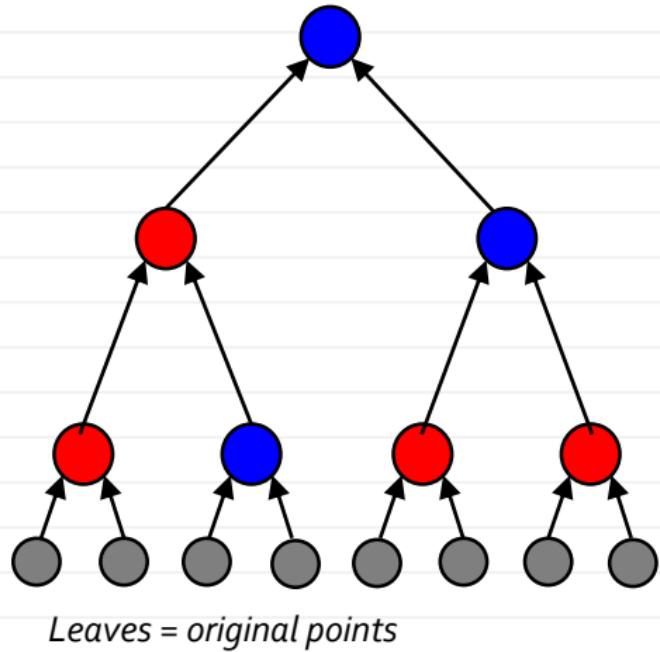
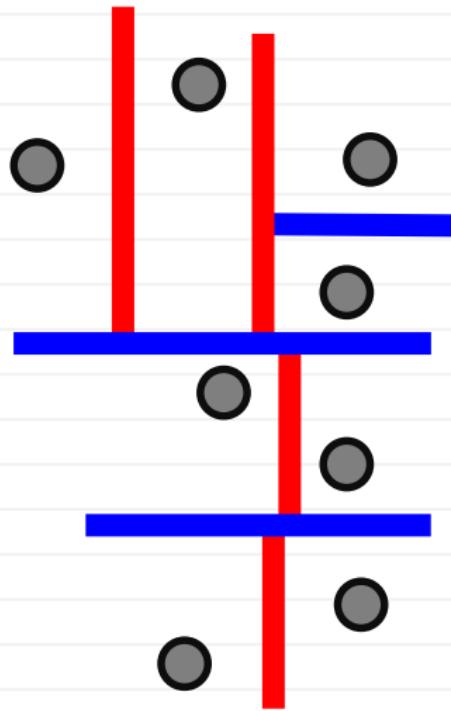
```
else return CreateLeaf(M[0])
```

Conventional use:

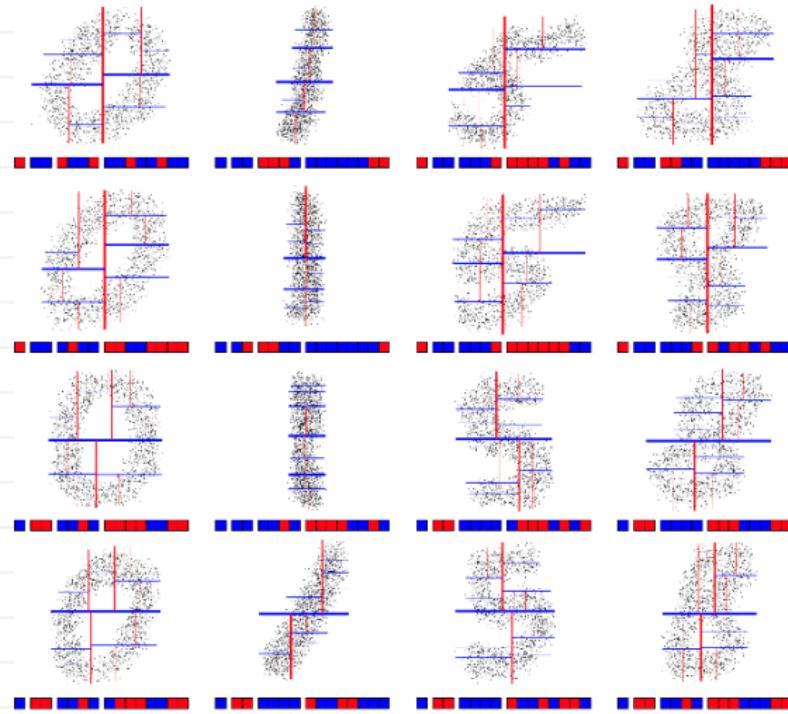
- Approximate nearest neighbor (alternate tree descent and backtracking)
- Ray tracing



## Kd-tree: 2D example

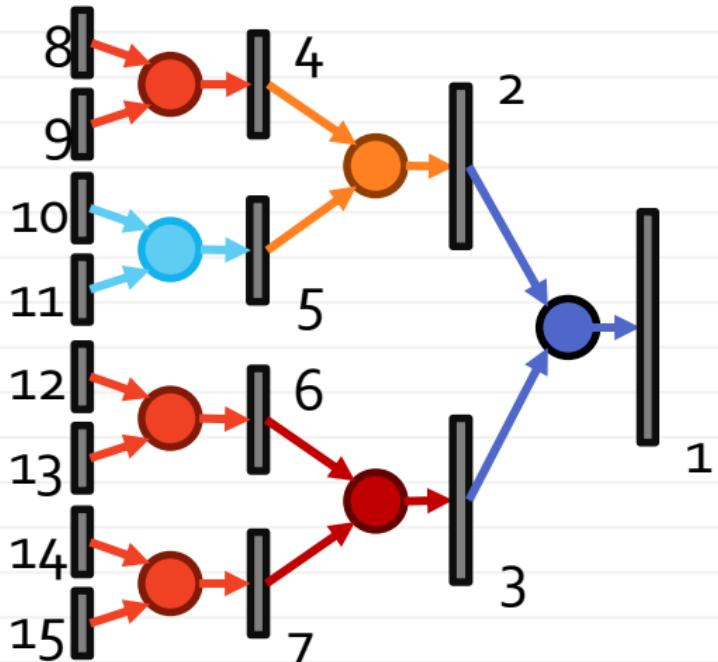
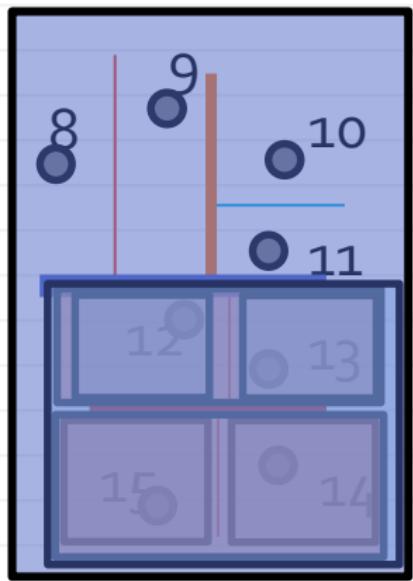


# Kd-tree for MNIST-dots



- 2D point clouds obtained from non-zero pixels in MNIST
- Kd-trees contain information about shapes
- Naïve kd-based classification: 82.4%

# Hierarchical features from a Kd-tree

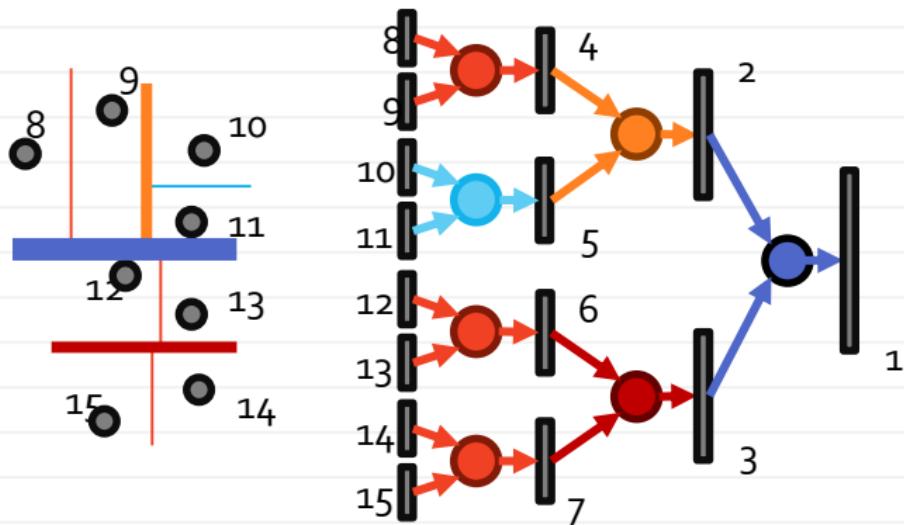


$$v_7 = \phi(W_7[v_{14}; v_{15}] + b_7)$$

$$v_6 = \phi(W_6[v_{12}; v_{13}] + b_6)$$

$$v_3 = \phi(W_3[v_6; v_7] + b_3)$$

# Hierarchical features from a Kd-tree



$$v_7 = \phi(W_7[v_{14}; v_{15}] + b_7)$$

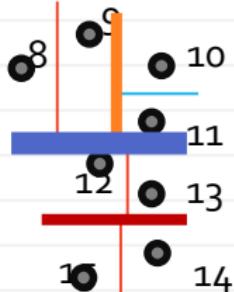


$$v_6 = \phi(W_6[v_{12}; v_{13}] + b_6)$$

# Kd-network “equation”

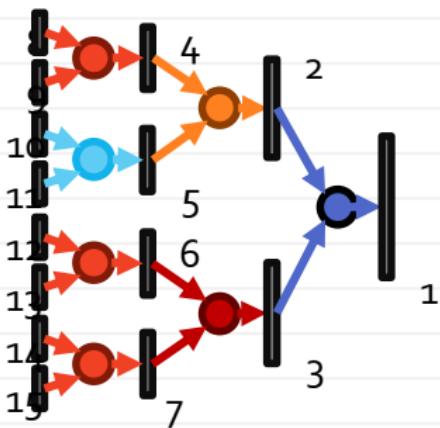
$$v_i = \begin{cases} \phi\left(W_x^{l_i}[v_{2i}; v_{2i+1}] + b_x^{l_i}\right) & \text{if } d_i = 'x' \\ \phi\left(W_y^{l_i}[v_{2i}; v_{2i+1}] + b_y^{l_i}\right) & \text{if } d_i = 'y' \\ \phi\left(W_z^{l_i}[v_{2i}; v_{2i+1}] + b_z^{l_i}\right) & \text{if } d_i = 'z' \end{cases}$$

$$v_i = \phi\left(W_{S_i}^{l_i}[v_{2i}; v_{2i+1}] + b_{S_i}^{l_i}\right)$$



ConvNet analogies:

- Localized linear operators
- Stationarity
- Bottom-up processing

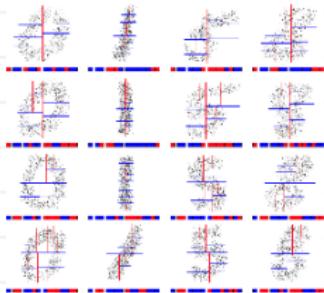


# Recap: Kd-network classifier

1. Build a Kd-tree
2. Run recursively bottom-up:

$$v_i = \phi \left( W_{S_i}^{l_i} [v_{2i}; v_{2i+1}] + b_{S_i}^{l_i} \right)$$

3. Output  $p(y|x) = \text{softmax}(W_0 v_1 + b_0)$

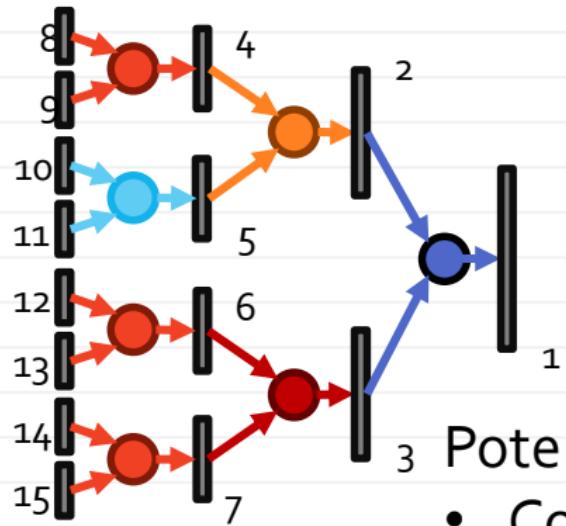


Learnable parameters

$$\left\{ \begin{array}{l} \{W_0; b_0\} \\ \{W_x^1; b_x^1; W_y^1; b_y^1; W_z^1; b_z^1\} \\ \{W_x^2; b_x^2; W_y^2; b_y^2; W_z^2; b_z^2\} \\ \{W_x^L; b_x^L; W_y^L; b_y^L; W_z^L; b_z^L\} \end{array} \right\}$$



# Kd-network: leaf representations



$[x, y, R, G, B]$



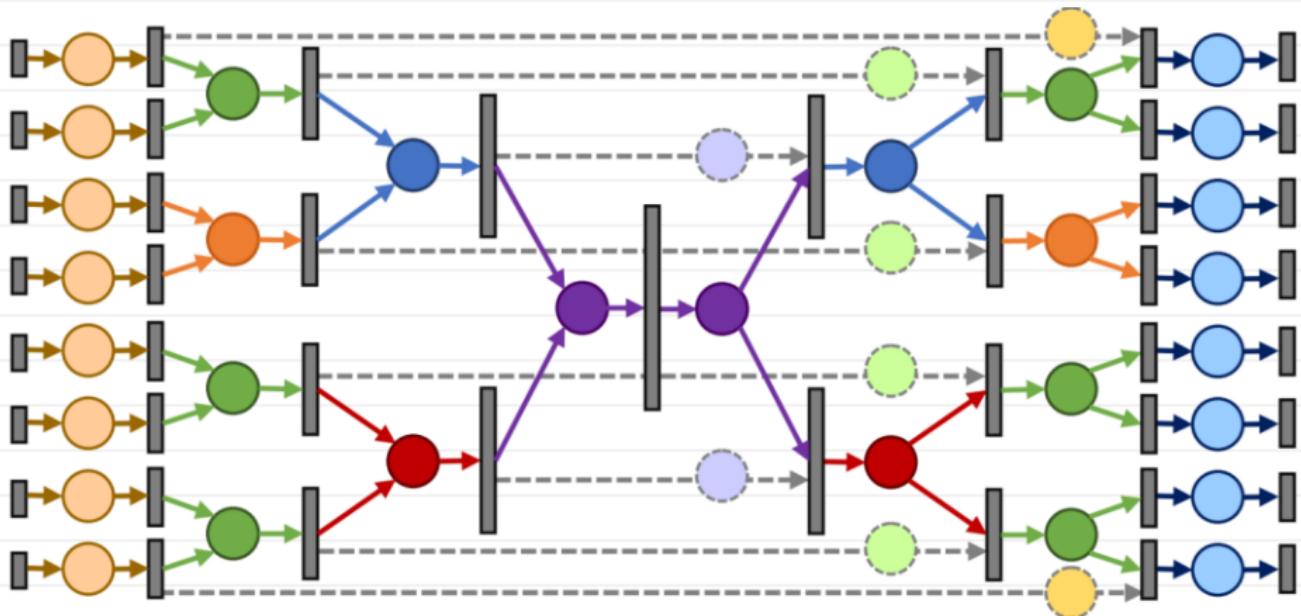
Potential options:

- Coordinates
- Color (R,G,B)
- Conv features
- Normal vector
- ..or nothing

[Klokov ICCV17]

# Kd-Nets for semantic segmentation

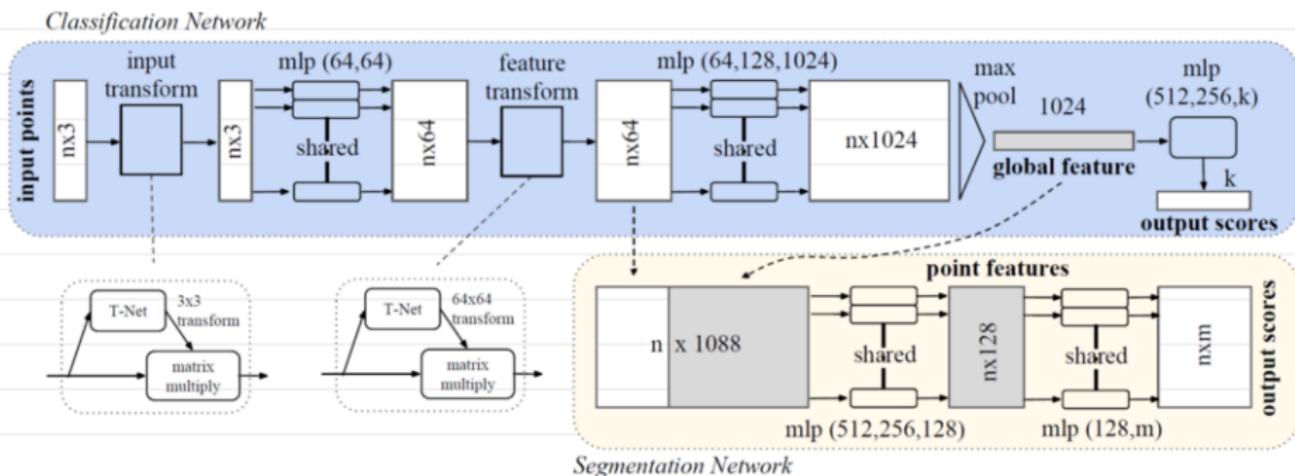
“Kd-Unet”:



[Klokov ICCV17]

# PointNet

Key idea: use only permutation invariant operators



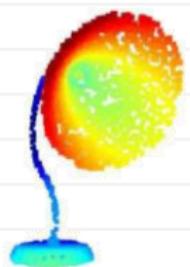
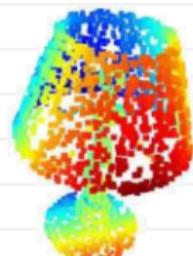
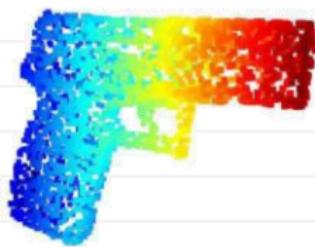
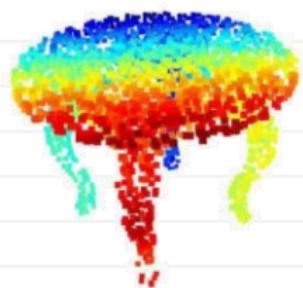
$T\text{-Net} = 1 \times 1 \text{ MLP} + \text{max-pool over points} + \text{MLP}$

Overall: efficient architecture with good performance

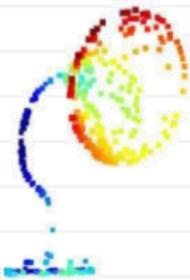
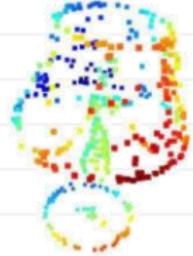
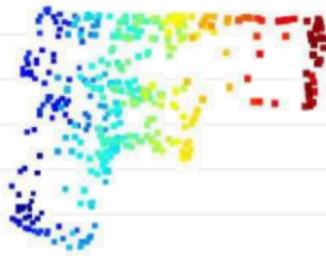
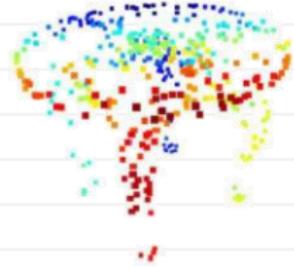
[Qi et al. CVPR17]

# PointNet results: critical sets

Original Shape



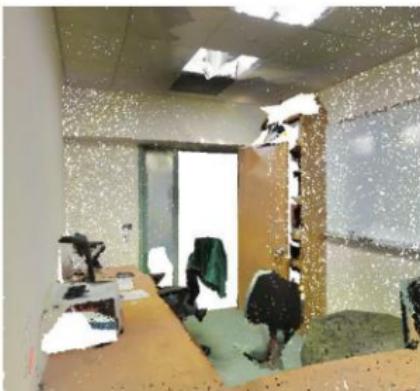
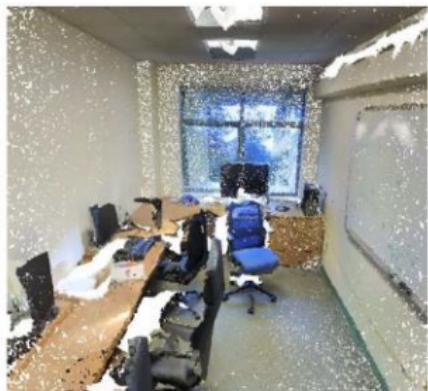
Critical Point Sets



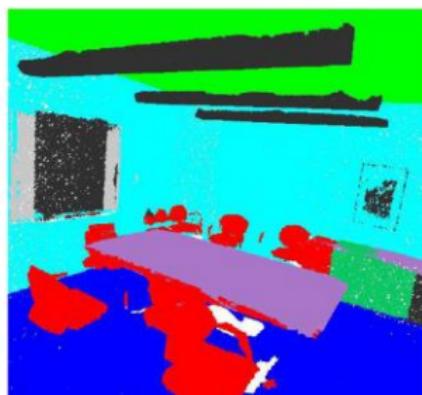
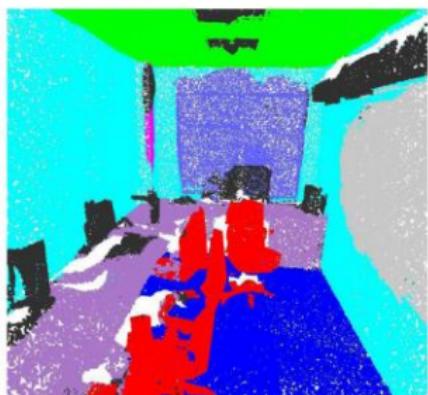
[Qi et al. CVPR17]

# PointNet results

Input

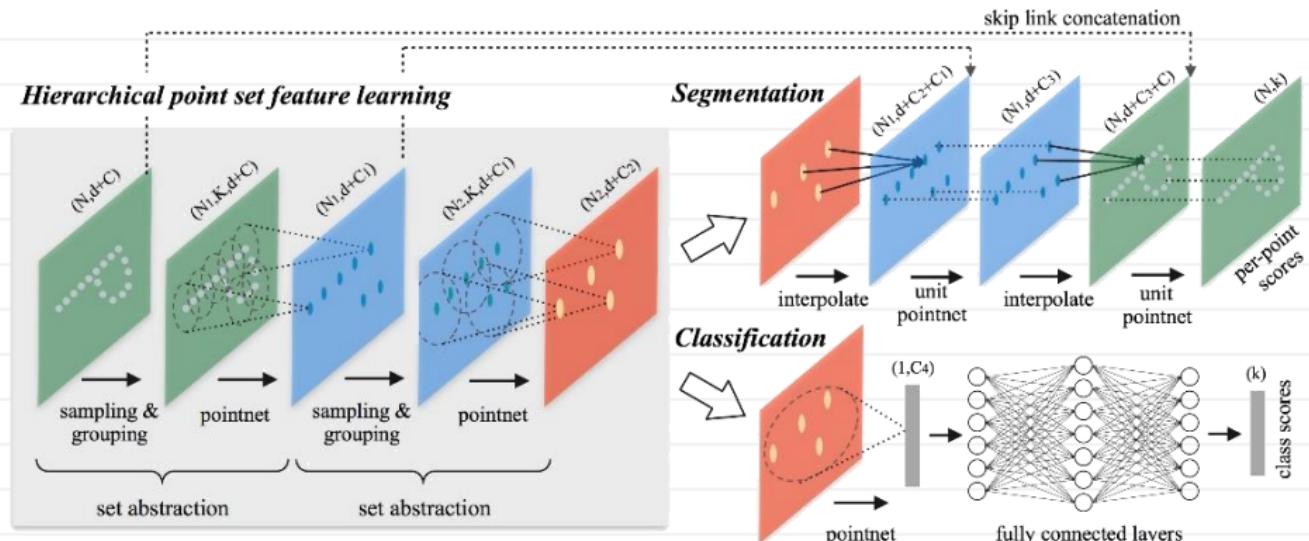


Output



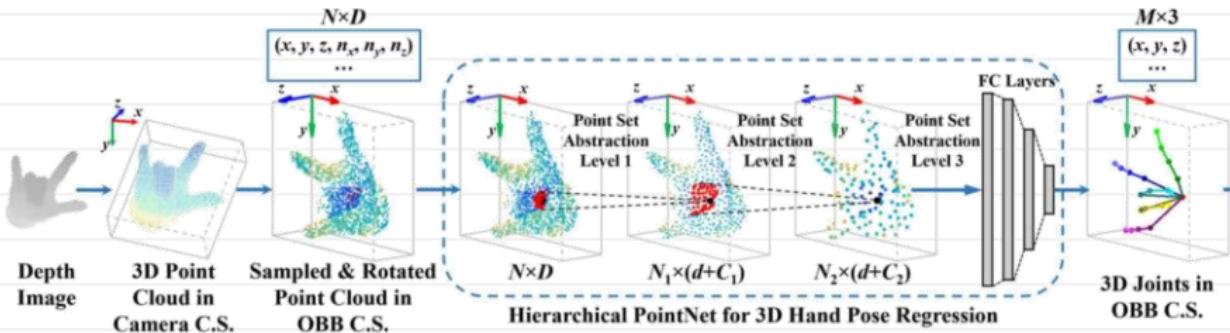
[Qi et al. CVPR17]

# PointNet++

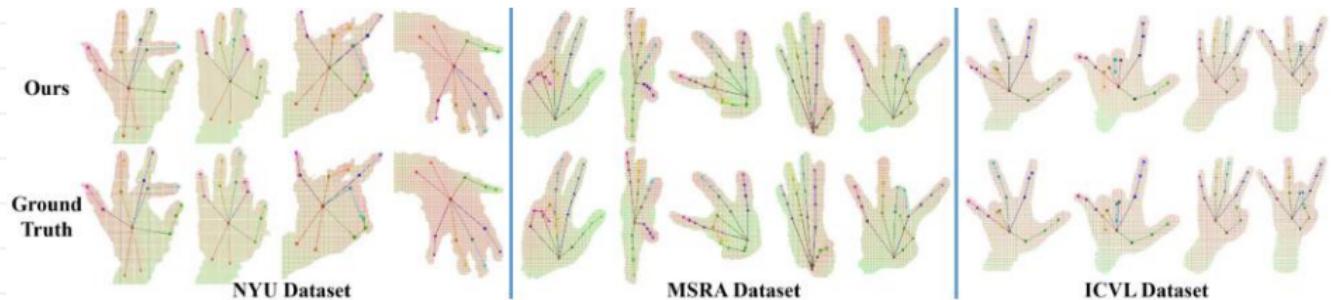


- Replaces global pooling with localized grouping and pooling
- Improved performance over PointNet (but slower)

# Hand PointNet

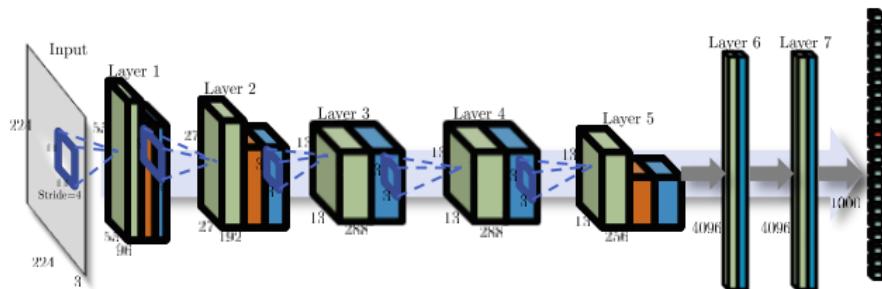


- Uses PointNet++ for hand keypoint estimation



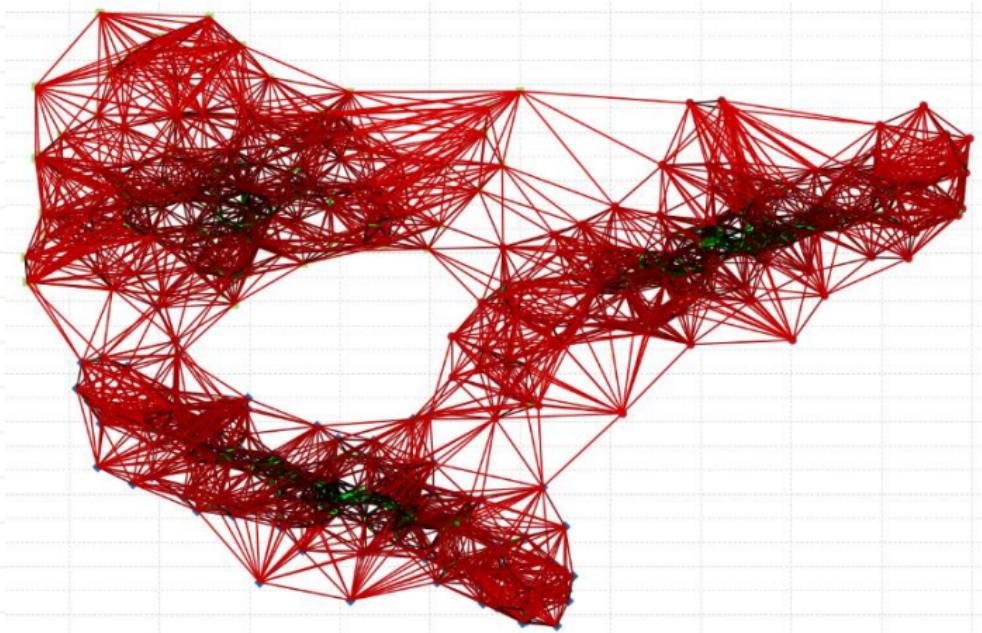
[Ge et al. CVPR18]

# General motivation



This is ConvNet  
It does recognition very well  
Be like ConvNet

# Point clouds as graphs



- k-nearest neighbor graphs
- hierarchical clustering graph
- ....

# Graph Laplacian

Graph Laplacian is a linear operator on graph function that measures the difference between the value at each vertex and adjacent neighbors

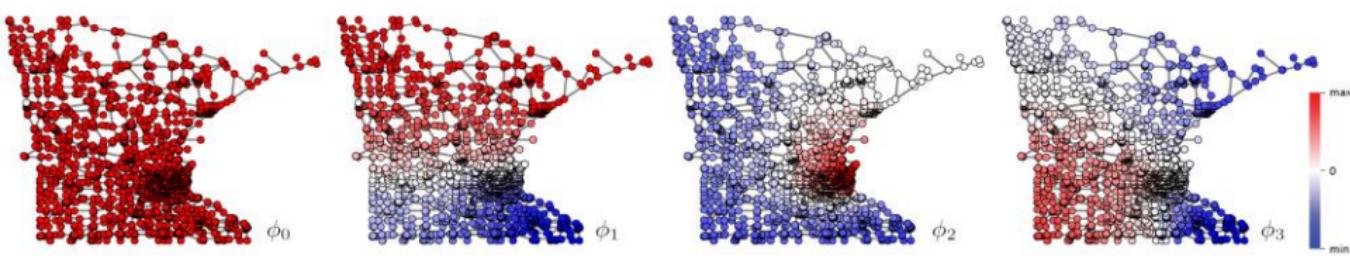
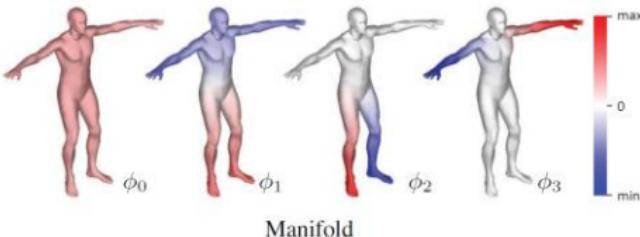
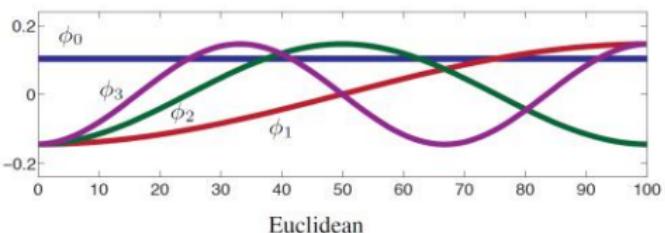
Labeled graph	Degree matrix	Adjacency matrix	Laplacian matrix
	$\begin{pmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$	$\begin{pmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{pmatrix}$

$$L_{\text{unnorm}} = D - A$$

$$L = I - D^{-1/2}AD^{-1/2}$$

- If  $x$  is a signal (function) on a graph, then  $Lx$  measures at vertex  $i$  measures the diff between  $x_i$  and the weighted average of the neighbors
- Laplacian measures pointwise smoothness

# Eigenbasis of Laplacian

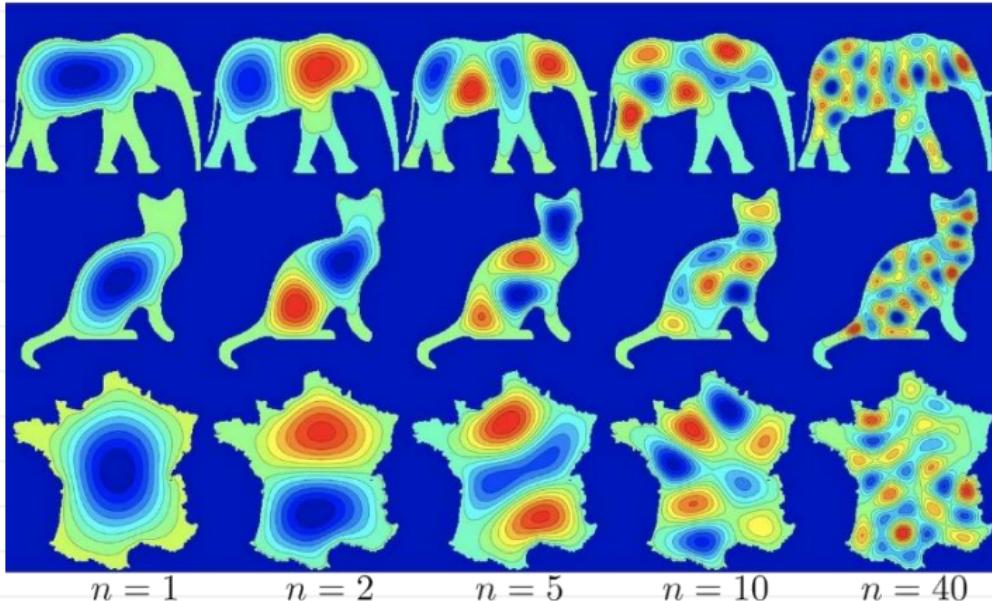


*Image source: [Bronstein et al. 2017]*

- For regular grids, eigenvectors are sine waves
- Projection of a signal onto the eigenbasis of the Laplacian generalizes Fourier transform to graphs

# Eigenbasis of Laplacian

Eigenvectors of the Laplacian:  $\Delta\psi_n = \lambda_n\psi_n$



*Image source:  
Gabriel Peyre*

- For regular grids, eigenvectors are sine functions
- Projection of a signal onto the eigenbasis of the Laplacian hence generalizes Fourier transform to graphs

# Graph-convolutional neural networks

- Projection of a signal on the eigenbasis of the Laplacian generalizes Fourier transform to graphs
- Kernel convolution is generalized as:

inverse graph Fourier transform

$$g_\theta \star x = U \operatorname{diag}(\theta) U^T x$$

Convolution parameters      graph Fourier transform

Neat, but:

- Not localized
- $O(n)$  learnable parameters
- Computationally heavy ( $U$  is dense)
- Poor generalization to new graphs

*Same story with standard convolutions (but we can impose restriction in the spatial domain)*

[Bruna ICLR14]

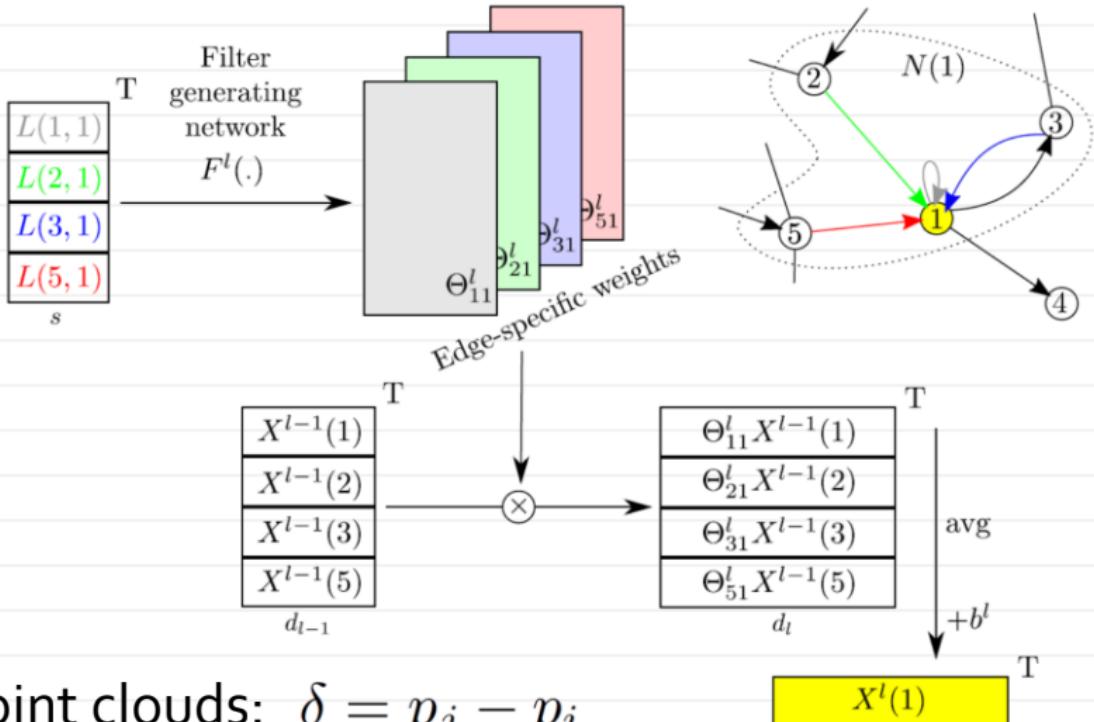
# Graph-convolutional neural networks

$$g_\theta \star x = U \operatorname{diag}(\theta) U^T x$$

- Laplacian is a convolution:  $Lx = U\Lambda U^T x$
- Degrees of the Laplacian are convolutions:  
$$L^k x = (U\Lambda U^T)^k x = U\Lambda^k U^T x$$
- Laplacian and its degrees are k-localized
- Idea: restrict considered filters to polynomials of degree K of the Laplacian (loose analogy: restricting spatial extent of filters in normal convolutions)
- The new operator is localized, has  $O(K)$  learnable parameters ( $K \ll n$ ), is efficient (no Fourier), generalizes reasonably well to new graphs
- Problem: not so expressive

[Defferrard et al. NIPS16]

# Dynamic edge conditioned filters

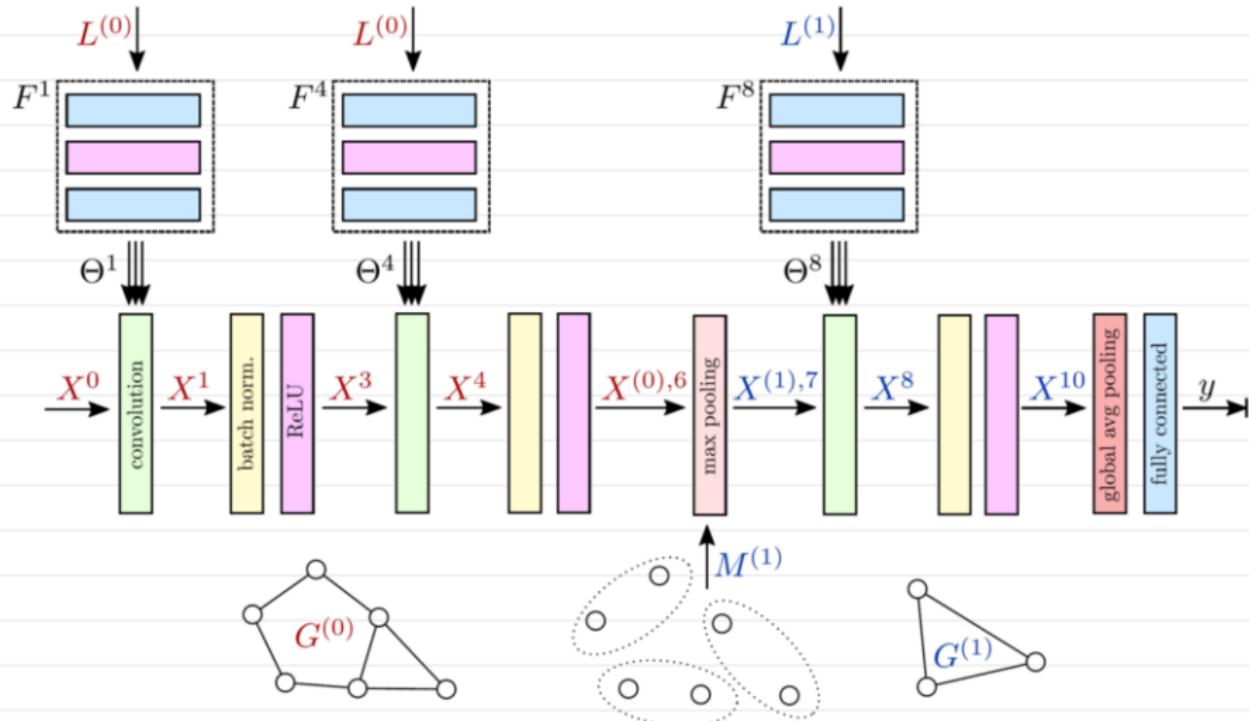


For point clouds:  $\delta = p_j - p_i$

$$L(j, i) = (\delta_x, \delta_y, \delta_z, \|\delta\|, \arccos \delta_z / \|\delta\|, \arctan \delta_y / \delta_x)$$

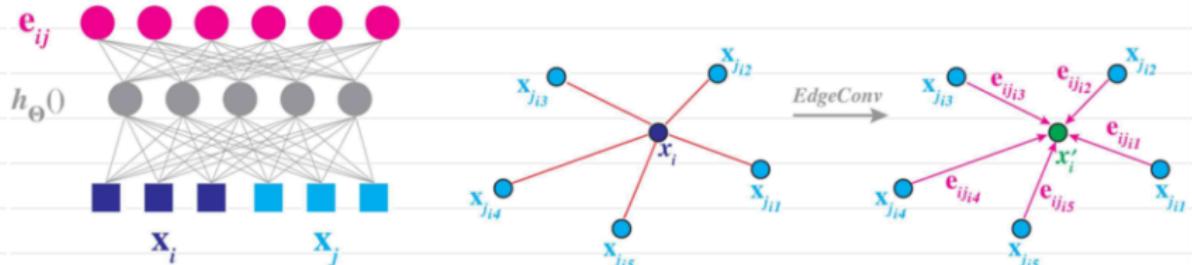
[Simonovsky CVPR17]

# Dynamic edge conditioned filters



[Simonovsky CVPR17]

# Dynamic Graph CNN



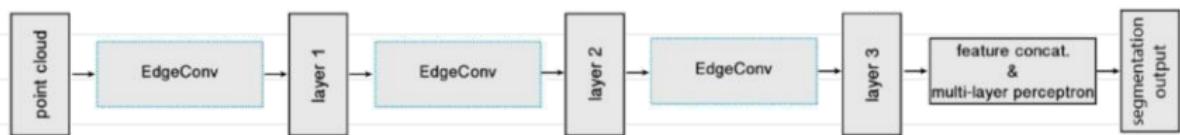
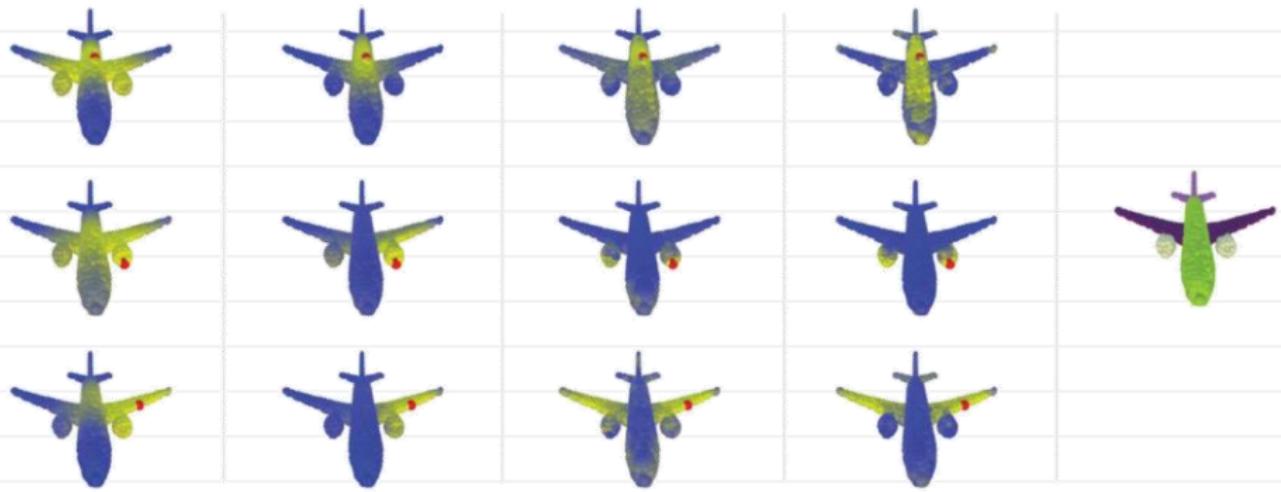
$$x'_i = \boxed{\quad} h_\Theta(x_i, x_j)$$

$\boxed{\quad}$  = sum or max

Very similar to [Simonovsky], but the graph is recomputed after each convolution

[Wang et al. CVPR18]

# Dynamic Graph CNN



[Wang et al. CVPR18]

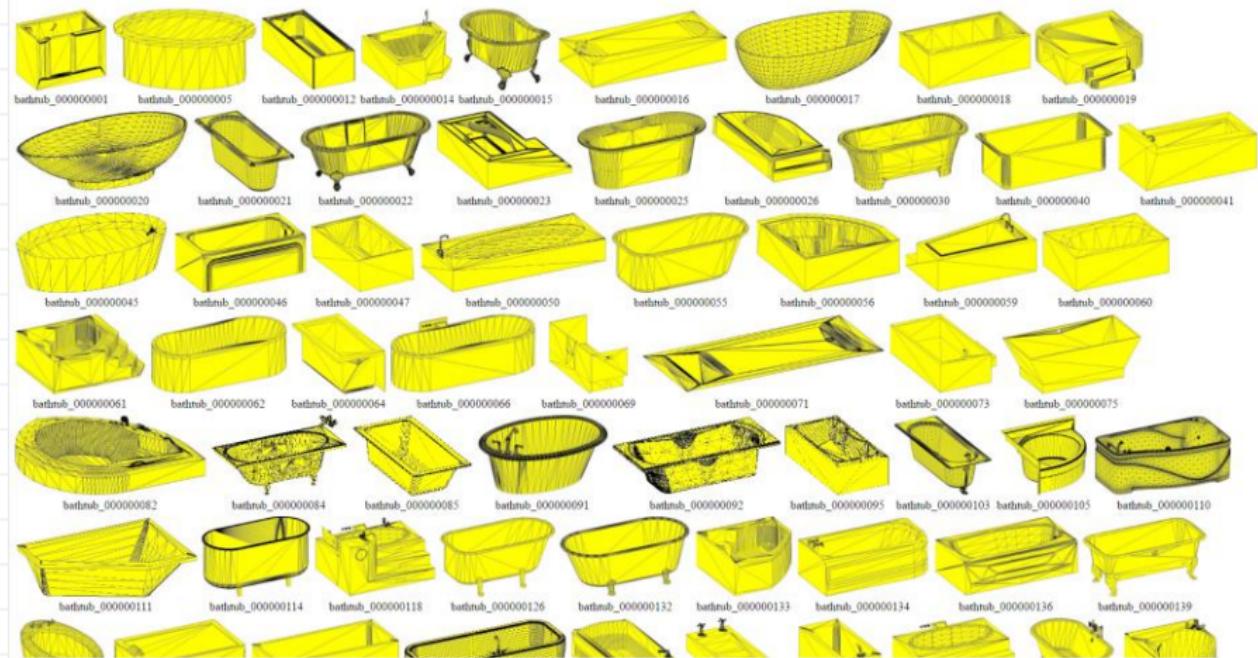
# Dynamic graph CNN

## ModelNet 40 benchmark:

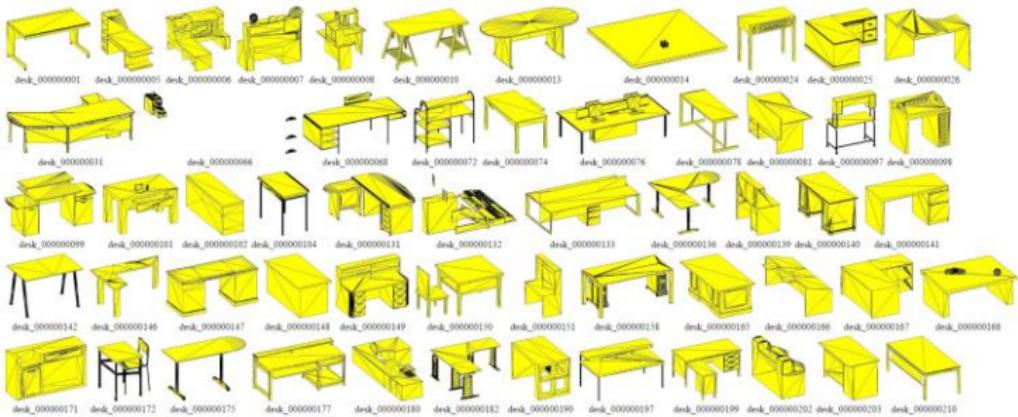
	MEAN CLASS ACCURACY	OVERALL ACCURACY
3D SHAPENETS [54]	77.3	84.7
VOXNET [30]	83.0	85.9
SUBVOLUME [35]	86.0	89.2
ECC [45]	83.2	87.4
POINTNET [34]	86.0	89.2
POINTNET++ [36]	-	90.7
KD-NET (DEPTH 10) [20]	-	90.6
KD-NET (DEPTH 15) [20]	-	91.8
OURS (BASELINE)	88.8	91.2
OURS	<b>90.2</b>	<b>92.2</b>

[Wang et al. CVPR18]

# ModelNet: bathtubs

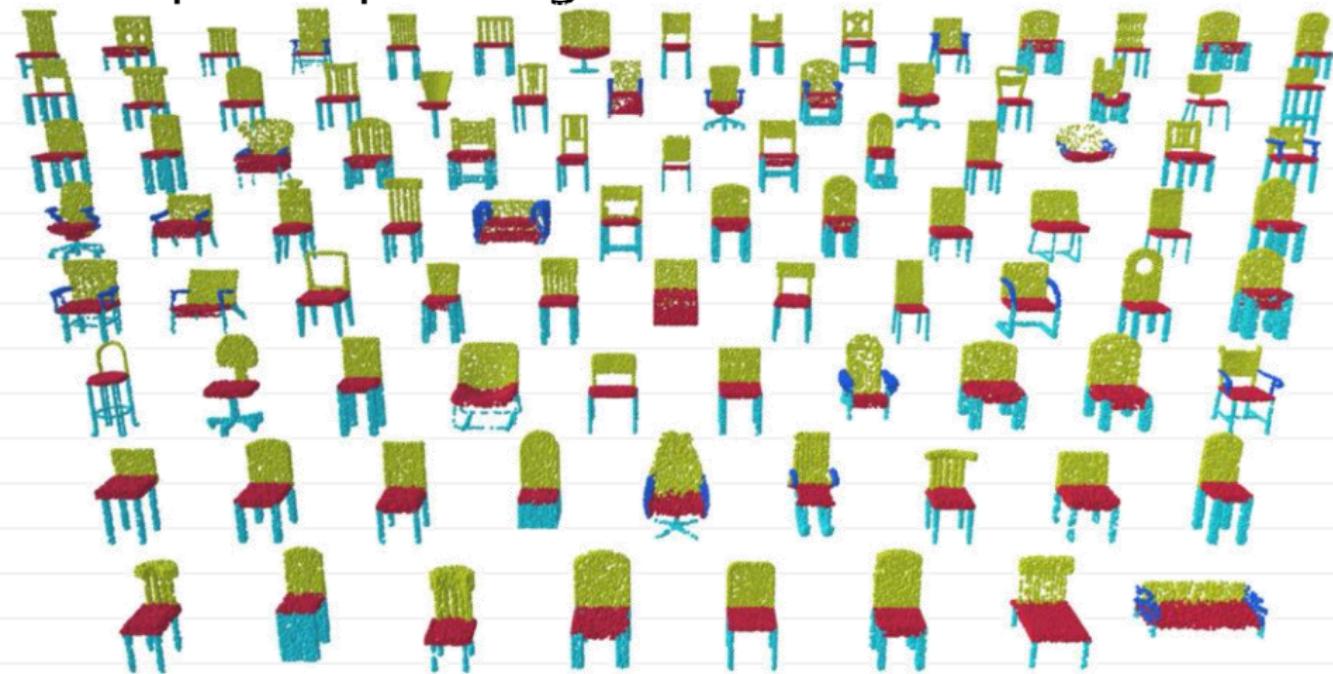


# ModelNet: desks and tables



# Dynamic graph CNN

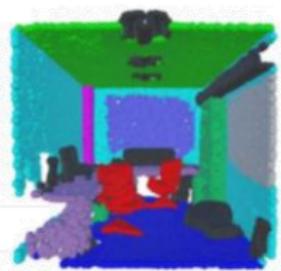
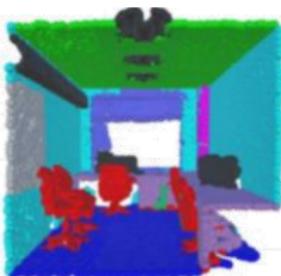
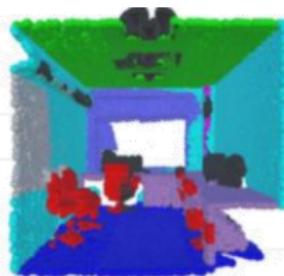
ShapeNet part segmentation benchmark:



[Wang et al. CVPR18]

# Dynamic graph CNN

Semantic segmentation results:



PointNet

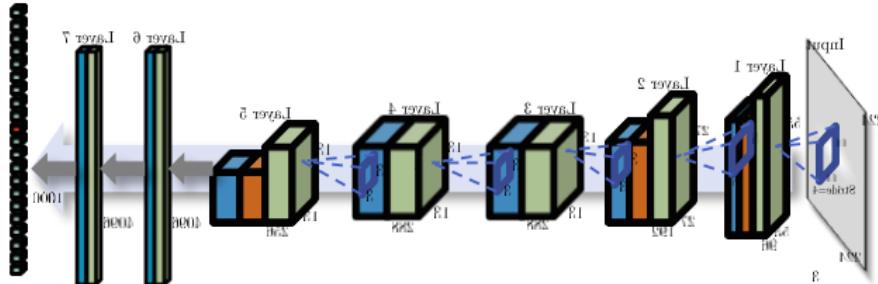
Ours

Ground truth

Real color

[Wang et al. CVPR18]

# General motivation

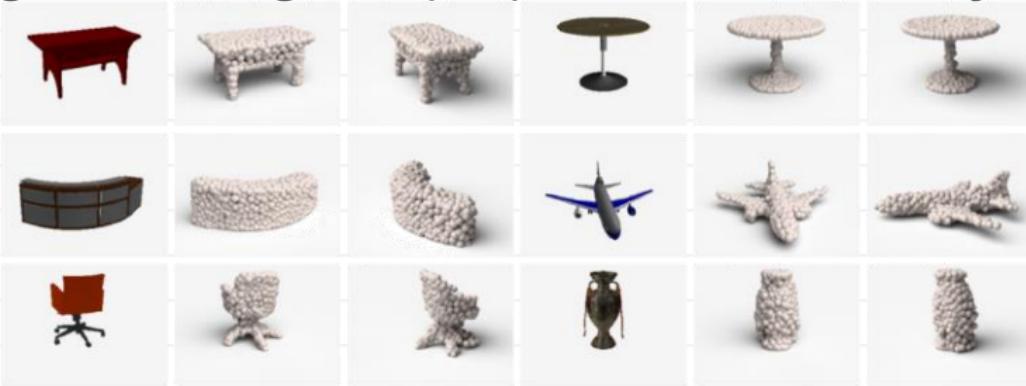


This is ConvNet  
It does generation very well  
Be like ConvNet

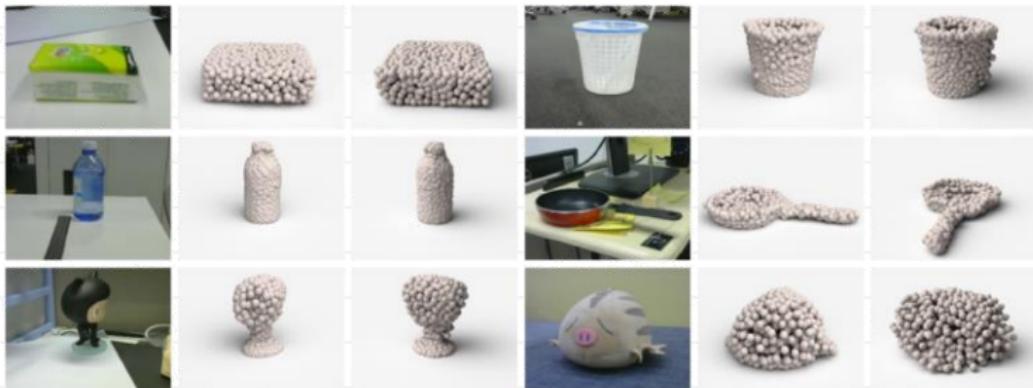
# Point Set Generation Networks

Task: given an image, output point cloud for the object:

Synthetic Data

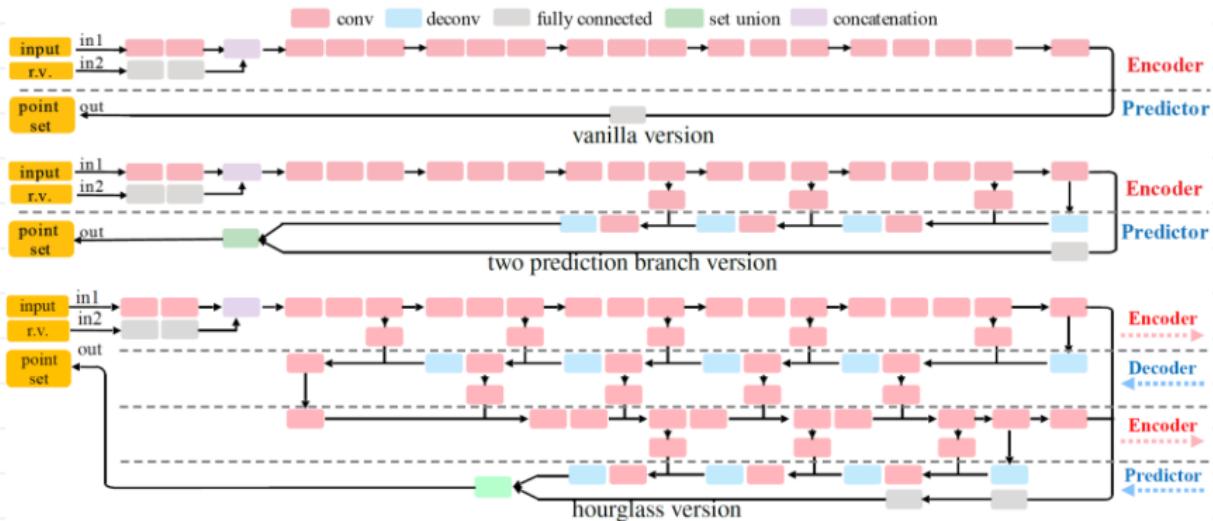


Real World Data



[Fan, Su, Guibas, CVPR17]

# Point Set Generation Networks



- Branch 1: raw point coordinates
- Branch 2: point coordinates packed into 2D XYZ maps

[Fan,Su,Guibas, CVPR17]

# Point Set Generation Networks



- Branch 1: raw point coordinates
- Branch 2: point coordinates packed into 2D XYZ maps

[Fan, Su, Guibas, CVPR17]

# Point Set Generation Networks: losses

Chamfer distance (CD):

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$

Earth-mover distance (EMD):

$$d_{EMD}(S_1, S_2) = \min_{\substack{\phi: S_1 \rightarrow S_2}} \sum_{x \in S_1} \|x - \phi(x)\|_2$$

*bijection*

[Fan, Su, Guibas, CVPR17]

# Point Set Generation Networks: losses



Chamfer distance (CD)

Earth-mover distance (EMD)

[Fan, Su, Guibas, CVPR17]

# Bibliography

- Charles Ruizhongtai Qi, Hao Su, Kaichun Mo, Leonidas J. Guibas:  
PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation. CVPR 2017: 77-85
- Charles Ruizhongtai Qi, Li Yi, Hao Su, Leonidas J. Guibas:  
PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space. NIPS 2017: 5105-5114
- Hang Su, Subhransu Maji, Evangelos Kalogerakis, Erik G. Learned-Miller:  
Multi-view Convolutional Neural Networks for 3D Shape Recognition. ICCV 2015: 945-953
- Hang Su, Varun Jampani, Deqing Sun, Subhransu Maji, Evangelos Kalogerakis, Ming-Hsuan Yang, Jan Kautz:  
SPLATNet: Sparse Lattice Networks for Point Cloud Processing. CVPR 2018: 2530-2539
- Daniel Maturana, Sebastian Scherer:  
VoxNet: A 3D Convolutional Neural Network for real-time object recognition. IROS 2015: 922-928
- Jonathan Tompson, Murphy Stein, Yann LeCun, Ken Perlin:  
Real-Time Continuous Pose Recovery of Human Hands Using Convolutional Networks. ACM Trans. Graph. 33(5): 169:1-169:10 (2014)
- Gyeongsik Moon, Ju Yong Chang, Kyoung Mu Lee:  
V2V-PoseNet: Voxel-to-Voxel Prediction Network for Accurate 3D Hand and Human Pose Estimation From a Single Depth Map. CVPR 2018: 5079-5088

# Bibliography

Roman Klokov, Victor S. Lempitsky:

Escape from Cells: Deep Kd-Networks for the Recognition of 3D Point Cloud Models. ICCV2017: 863-872

Benjamin Graham, Martin Engelcke, Laurens van der Maaten:

3D Semantic Segmentation With Submanifold Sparse Convolutional Networks. CVPR2018: 9224-9232

Liuhan Ge, Yujun Cai, Junwu Weng, Junsong Yuan:

Hand PointNet: 3D Hand Pose Estimation Using Point Sets. CVPR 2018: 8417-8426

Michael M. Bronstein, Joan Bruna, Yann LeCun, Arthur Szlam, Pierre Vandergheynst:

Geometric Deep Learning: Going beyond Euclidean data. IEEE Signal Process. Mag. 34(4): 18-42 (2017)

Joan Bruna, Wojciech Zaremba, Arthur Szlam, Yann LeCun:

Spectral Networks and Locally Connected Networks on Graphs. ICLR 2014

Michaël Defferrard, Xavier Bresson, Pierre Vandergheynst:

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. NIPS2016: 3837-3845

Yue Wang, Yongbin Sun, Ziwei Liu, Sanjay E. Sarma, Michael M. Bronstein, Justin M. Solomon:

Dynamic Graph CNN for Learning on Point Clouds. CoRR abs/1801.07829 (2018)

Martin Simonovsky, Nikos Komodakis:

Dynamic Edge-Conditioned Filters in Convolutional Neural Networks on Graphs. CVPR 2017: 29-38

Haoqiang Fan, Hao Su, Leonidas J. Guibas:

A Point Set Generation Network for 3D Object Reconstruction from a Single Image. CVPR 2017: 2463-2471